

# PARTIE 4 : Observabilité avec Prometheus et Grafana

Durée : 45 minutes

## Objectifs

Dans cette quatrième partie, vous allez déployer une stack complète d'observabilité avec Prometheus et Grafana pour moniter la plateforme CloudShop.

**Compétences évaluées :** - Installation de Prometheus Operator - Configuration des ServiceMonitors pour scraping - Création de dashboards Grafana - Définition d'alerting rules - Calcul et monitoring des SLO/SLI

---

## Travail à Réaliser

### Tâche 4.1 : Installation de Prometheus et Grafana (4 points)

*Installation via Kube-Prometheus-Stack*

*# 1. Ajouter le repo Helm*

```
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
```

```
helm repo update
```

*# 2. Créer le namespace monitoring*

```
kubectl create namespace monitoring
```

*# 3. Installer la stack complète (Prometheus + Grafana + Alertmanager)*

```
helm install prometheus prometheus-community/kube-prometheus-stack \
--namespace monitoring \
--set prometheus.prometheusSpec.serviceMonitorSelectorNilUsesHelmValues=false \
--set prometheus.prometheusSpec.podMonitorSelectorNilUsesHelmValues=false \
--set grafana.adminPassword=admin123
```

*# 4. Attendre que tous les pods soient prêts*

```
kubectl wait --for=condition=ready pod --all -n monitoring --timeout=300s
```

*# 5. Vérifier l'installation*

```
kubectl get pods -n monitoring  
kubectl get svc -n monitoring
```

### *Accès aux Interfaces*

#### *# Prometheus*

```
kubectl port-forward svc/prometheus-operated 9090:9090 -n monitoring  
# URL : http://localhost:9090
```

#### *# Grafana*

```
kubectl port-forward svc/prometheus-grafana 3000:80 -n monitoring  
# URL : http://localhost:3000  
# User : admin  
# Password : admin123
```

#### *# Alertmanager*

```
kubectl port-forward svc/prometheus-kube-prometheus-alertmanager 9093:9093 -n monitoring  
# URL : http://localhost:9093
```

---

## Tâche 4.2 : ServiceMonitors pour Scraping (4 points)

Les ServiceMonitors configurent Prometheus pour scraper les métriques des microservices.

### *ServiceMonitor pour API Gateway*

**Fichier** : monitoring/servicemonitors/api-gateway.yaml

```
apiVersion: monitoring.coreos.com/v1  
kind: ServiceMonitor  
metadata:  
  name: api-gateway  
  namespace: monitoring  
  labels:  
    app: api-gateway  
    release: prometheus  
spec:  
  selector:  
    matchLabels:  
      app: api-gateway  
  namespaceSelector:  
    matchNames:
```

```
- cloudshop-prod  
endpoints:  
- port: http  
  interval: 30s  
  path: /metrics  
  scheme: http
```

### *ServiceMonitor pour Auth Service*

**Fichier** : monitoring/servicemonitors/auth-service.yaml

```
apiVersion: monitoring.coreos.com/v1  
kind: ServiceMonitor  
metadata:  
  name: auth-service  
  namespace: monitoring  
  labels:  
    app: auth-service  
    release: prometheus  
spec:  
  selector:  
    matchLabels:  
      app: auth-service  
  namespaceSelector:  
    matchNames:  
    - cloudshop-prod  
  endpoints:  
  - port: http  
    interval: 30s  
    path: /metrics  
    scheme: http
```

### *ServiceMonitor pour Products API*

**Fichier** : monitoring/servicemonitors/products-api.yaml

```
apiVersion: monitoring.coreos.com/v1  
kind: ServiceMonitor  
metadata:  
  name: products-api  
  namespace: monitoring  
  labels:  
    app: products-api  
    release: prometheus  
spec:
```

```
selector:  
  matchLabels:  
    app: products-api  
namespaceSelector:  
  matchNames:  
    - cloudshop-prod  
endpoints:  
  - port: http  
    interval: 30s  
    path: /metrics  
    scheme: http
```

### *ServiceMonitor pour Orders API*

**Fichier** : monitoring/servicemonitors/orders-api.yaml

```
apiVersion: monitoring.coreos.com/v1  
kind: ServiceMonitor  
metadata:  
  name: orders-api  
  namespace: monitoring  
  labels:  
    app: orders-api  
    release: prometheus  
spec:  
  selector:  
    matchLabels:  
      app: orders-api  
  namespaceSelector:  
    matchNames:  
      - cloudshop-prod  
  endpoints:  
    - port: http  
      interval: 30s  
      path: /metrics  
      scheme: http
```

### *Déploiement des ServiceMonitors*

*# Appliquer tous les ServiceMonitors*

```
kubectl apply -f monitoring/servicemonitors/
```

*# Vérifier*

```
kubectl get servicemonitor -n monitoring
```

```
# Vérifier que Prometheus les détecte
# Aller dans Prometheus UI > Status > Targets
# Les targets cloudshop-prod/* doivent apparaître
```

---

## Tâche 4.3 : Alerting Rules (3 points)

Définissez des règles d'alerte pour détecter les anomalies.

**Fichier** : monitoring/alerts/prometheus-rules.yaml

```
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  name: cloudshop-alerts
  namespace: monitoring
  labels:
    release: prometheus
spec:
  groups:
  - name: cloudshop.rules
    interval: 30s
    rules:
      # Alert : Taux d'erreur élevé (>5%)
      - alert: HighErrorRate
        expr: |
          (
            sum(rate(http_requests_total{status=~"5..",namespace="cloudshop-prod"}[5m])) /
            sum(rate(http_requests_total{namespace="cloudshop-prod"}[5m]))
          ) > 0.05
        for: 5m
        labels:
          severity: warning
          component: backend
        annotations:
          summary: "Taux d'erreur élevé détecté"
          description: "Le service {{ $labels.service }} a un taux d'erreur de {{ $value | humanizePercentage }} sur les 5 dernières minutes."
      # Alert : Latence P95 élevée (>1s)
      - alert: HighLatencyP95
```

```
expr: |
  histogram_quantile(0.95,
    sum(rate(http_request_duration_seconds_bucket{namespace="cloudshop-prod"})
[5m])) by (le, service)
  ) > 1
for: 10m
labels:
  severity: warning
  component: performance
annotations:
  summary: "Latence P95 élevée"
  description: "Le service {{ $labels.service }} a une latence P95 de {{ $value }}s."
# Alerte : Pod CrashLooping
- alert: PodCrashLooping
expr: |
  rate(kube_pod_container_status_restarts_total{namespace="cloudshop-prod"}[15m])
> 0
for: 5m
labels:
  severity: critical
  component: infrastructure
annotations:
  summary: "Pod en CrashLoop détecté"
  description: "Le pod {{ $labels.pod }} redémarre fréquemment."
# Alerte : Pod non disponible
- alert: PodNotReady
expr: |
  kube_pod_status_phase{namespace="cloudshop-prod",phase!="Running"} > 0
for: 5m
labels:
  severity: warning
  component: infrastructure
annotations:
  summary: "Pod non disponible"
  description: "Le pod {{ $labels.pod }} n'est pas en état Running."
# Alerte : Utilisation mémoire élevée (>80%)
- alert: HighMemoryUsage
expr: |
(
  container_memory_working_set_bytes{namespace="cloudshop-prod"}
```

```

/
  container_spec_memory_limit_bytes{namespace="cloudshop-prod"}
) > 0.8
for: 10m
labels:
  severity: warning
  component: resources
annotations:
  summary: "Utilisation mémoire élevée"
  description: "Le conteneur {{ $labels.container }} utilise {{ $value | humanizePercentage }} de sa limite mémoire."
# Alerte : SLO Breach (Availability < 99.9%)
- alert: SLOBreach
  expr: |
    (
      sum(rate(http_requests_total{status!~"5..",namespace="cloudshop-prod"}[30m])) /
      sum(rate(http_requests_total{namespace="cloudshop-prod"}[30m]))
    ) < 0.999
  for: 5m
  labels:
    severity: critical
    component: slo
  annotations:
    summary: "SLO de disponibilité non respecté"
    description: "La disponibilité est de {{ $value | humanizePercentage }}, en dessous du SLO de 99.9%."
# Appliquer les règles
kubectl apply -f monitoring/alerts/prometheus-rules.yaml

# Vérifier dans Prometheus UI
# Status > Rules
# Alerts

```

---

## Tâche 4.4 : Dashboards Grafana (4 points)

*Dashboard 1 : CloudShop Overview*

**Fichier :** monitoring/dashboards/cloudshop-overview.json

```
{
  "dashboard": {
    "title": "CloudShop - Overview",
    "panels": [
      {
        "title": "Request Rate (RPS)",
        "targets": [
          {
            "expr": "sum(rate(http_requests_total{namespace=\"cloudshop-prod\"}[5m]))",
            "legendFormat": "Total RPS"
          }
        ],
        "type": "graph"
      },
      {
        "title": "Error Rate (%)",
        "targets": [
          {
            "expr": "sum(rate(http_requests_total{status=~\"5..\",namespace=\"cloudshop-prod\"}[5m])) / sum(rate(http_requests_total{namespace=\"cloudshop-prod\"}[5m])) * 100",
            "legendFormat": "Error Rate"
          }
        ],
        "type": "graph"
      },
      {
        "title": "P95 Latency",
        "targets": [
          {
            "expr": "histogram_quantile(0.95,sum(rate(http_request_duration_seconds_bucket{namespace=\"cloudshop-prod\"}[5m])) by (le))",
            "legendFormat": "P95 Latency"
          }
        ],
        "type": "graph"
      },
      {
        "title": "Pods Status",
        "targets": [
          {
            "expr": "count(kube_pod_status_phase{namespace=\"cloudshop-"
        }
      }
    ]
  }
}
```

```

    prod\",phase=\"Running\"}",
      "legendFormat": "Running"
    },
  {
    "expr": "count(kube_pod_status_phase{namespace=\"cloudshop-prod\",phase!=\"Running\"})",
      "legendFormat": "Not Running"
    }
  ],
  "type": "stat"
}
]
}
}

```

## *Dashboard 2 : SLO Monitoring*

### **Queries PromQL pour le dashboard SLO :**

```

# Availability SLO (99.9%)
(
  sum(rate(http_requests_total{status=~"5..",namespace="cloudshop-prod"}[30d]))
  /
  sum(rate(http_requests_total{namespace="cloudshop-prod"}[30d]))
) * 100

# Error Budget Remaining (30 days)
(
  1 - (
    sum(rate(http_requests_total{status=~"5..",namespace="cloudshop-prod"}[30d]))
    /
    sum(rate(http_requests_total{namespace="cloudshop-prod"}[30d]))
  )
  - 0.999
) / (1 - 0.999) * 100

# Burn Rate (1h)
(
  sum(rate(http_requests_total{status=~"5..",namespace="cloudshop-prod"}[1h]))
  /
  sum(rate(http_requests_total{namespace="cloudshop-prod"}[1h]))
) / (1 - 0.999)

```

```
# Latency SLI (P95 < 500ms)
histogram_quantile(0.95,
  sum(rate(http_request_duration_seconds_bucket{namespace="cloudshop-prod"}[5m]))
by (le)
) * 1000
```

### *Import des Dashboards dans Grafana*

**Méthode 1 : Via UI** 1. Grafana UI > Dashboards > Import 2. Copier/coller le JSON  
3. Sélectionner Prometheus datasource 4. Save

### **Méthode 2 : Via ConfigMap**

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: grafana-dashboards
  namespace: monitoring
  labels:
    grafana_dashboard: "1"
data:
  cloudshop-overview.json: |
    # Contenu du JSON dashboard
```

```
kubectl apply -f monitoring/dashboards/configmap-dashboards.yaml
```

---

## Définition des SLO/SLI

### SLO (Service Level Objective)

**Availability SLO** : 99.9% uptime sur 30 jours glissants

**Calcul** : - Downtime autorisé =  $(1 - 0.999) * 30 \text{ jours} * 24\text{h} * 60\text{min} = 43.2 \text{ minutes/mois}$

**Error Budget** : - Budget initial = 43.2 minutes - Budget consommé = downtime réel - Budget restant = 43.2 - downtime réel

### SLI (Service Level Indicator)

**Availability SLI** : Ratio de requêtes réussies

```
sum(rate(http_requests_total{status!="5.."}[5m]))  
/  
sum(rate(http_requests_total[5m]))
```

**Latency SLI** : P95 latency < 500ms

```
histogram_quantile(0.95,  
  sum(rate(http_request_duration_seconds_bucket[5m])) by (le)  
) < 0.5
```

---

## Validation

### Checklist

# 1. Stack Prometheus/Grafana installée

```
kubectl get pods -n monitoring
```

```
# prometheus-prometheus-kube-prometheus-prometheus-0  Running
```

```
# prometheus-grafana-xxx           Running
```

```
# prometheus-kube-prometheus-operator-xxx     Running
```

```
# alertmanager-prometheus-kube-prometheus-alertmanager-0  Running
```

# 2. ServiceMonitors créés

```
kubectl get servicemonitor -n monitoring
```

```
# Doit lister : api-gateway, auth-service, products-api, orders-api
```

# 3. Prometheus scrape les targets

```
# Ouvrir http://localhost:9090 > Status > Targets
```

```
# Les targets cloudshop-prod/* doivent être UP
```

# 4. Alerting rules configurées

```
# Prometheus UI > Alerts
```

```
# Les alertes cloudshop.rules doivent apparaître
```

# 5. Grafana accessible avec dashboards

```
# Ouvrir http://localhost:3000
```

```
# Dashboards > CloudShop Overview et SLO Monitoring doivent exister
```

# 6. Métriques affichées

```
# Dans Grafana, les panels doivent afficher des données (pas vides)
```

## Queries de Test dans Prometheus

```
# 1. Total requests per second
sum(rate(http_requests_total{namespace="cloudshop-prod"}[5m]))
```

  

```
# 2. Error rate
sum(rate(http_requests_total{status=~"5..",namespace="cloudshop-prod"}[5m])) /
sum(rate(http_requests_total{namespace="cloudshop-prod"}[5m]))
```

  

```
# 3. P95 latency
histogram_quantile(0.95,
sum(rate(http_request_duration_seconds_bucket{namespace="cloudshop-prod"}[5m])) by (le))
```

  

```
# 4. Pods running
count(kube_pod_status_phase{namespace="cloudshop-prod",phase="Running"})
```

  

```
# 5. CPU usage by pod
sum(rate(container_cpu_usage_seconds_total{namespace="cloudshop-prod"}[5m])) by (pod)
```

  

```
# 6. Memory usage by pod
sum(container_memory_working_set_bytes{namespace="cloudshop-prod"}) by (pod) /
1024 / 1024
```

---

## ServiceMonitor ne scrape pas

# Vérifier les labels

```
kubectl get servicemonitor <name> -n monitoring -o yaml
```

# Vérifier le selector dans Prometheus

```
kubectl get prometheus -n monitoring -o yaml | grep serviceMonitorSelector
```

# Solution : Ajouter le label release: prometheus

```
metadata:
```

```
labels:
```

```
release: prometheus
```

## Aucune métrique dans Grafana

# 1. Vérifier Prometheus datasource dans Grafana

# Configuration > Data sources > Prometheus

```
# URL : http://prometheus-operated:9090
```

```
# 2. Tester une query simple  
# Explore > Prometheus > up
```

```
# 3. Vérifier que les targets sont UP  
# Prometheus UI > Status > Targets
```

Alerting rules ne s'affichent pas

```
# Vérifier les labels de la PrometheusRule  
kubectl get prometheusrule -n monitoring -o yaml
```

```
# Doit avoir le label : release: prometheus
```

```
# Vérifier que Prometheus les charge  
# Prometheus UI > Status > Rules
```

---

## Ressources

- [Prometheus Operator](#)
  - [PromQL Basics](#)
  - [Grafana Dashboards](#)
  - [SLO Best Practices](#)
- 

**Une fois cette partie terminée, passez à PARTIE5\_SECURITE\_SRE.md**